

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv("recipe_site_traffic_2212.csv")
```

```
In [3]: df.head()
```

```
Out[3]:   recipe  calories  carbohydrate  sugar  protein  category  servings  high_traffic
0         1        NaN           NaN     NaN       NaN      Pork        6      High
1         2      35.48        38.56    0.66      0.92    Potato        4      High
2         3     914.28        42.68    3.09      2.88  Breakfast        1      NaN
3         4      97.03        30.56    38.63     0.02  Beverages        4      High
4         5      27.05        1.85     0.80      0.53  Beverages        4      NaN
```



```
In [4]: df.isna().sum()
```

```
Out[4]:   recipe      0
calories     52
carbohydrate  52
sugar        52
protein      52
category      0
servings      0
high_traffic  373
dtype: int64
```

```
In [5]: for i in df.columns:
    print(df[i].dtypes)
```

```
int64
float64
float64
float64
float64
object
object
object
```

```
In [6]: df["high_traffic"].value_counts()
```

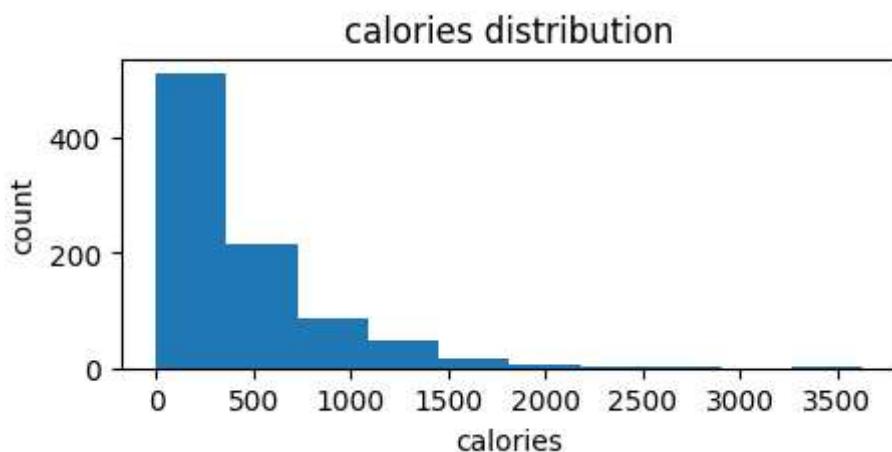
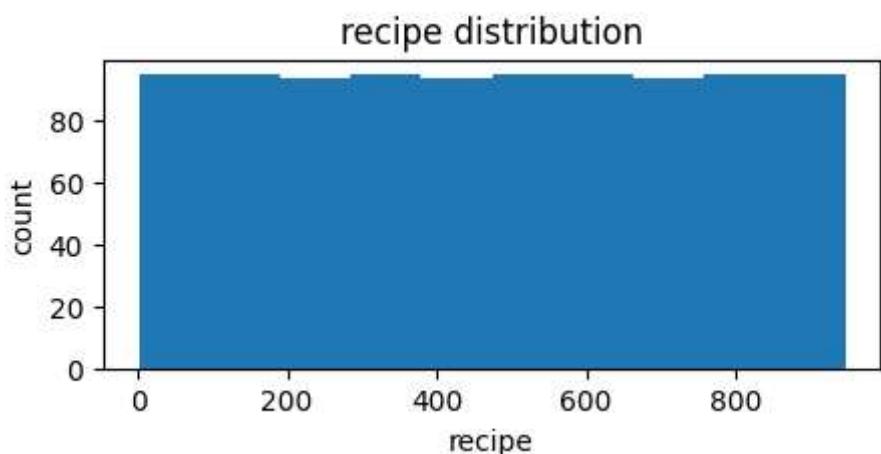
```
Out[6]: high_traffic
High      574
Name: count, dtype: int64
```

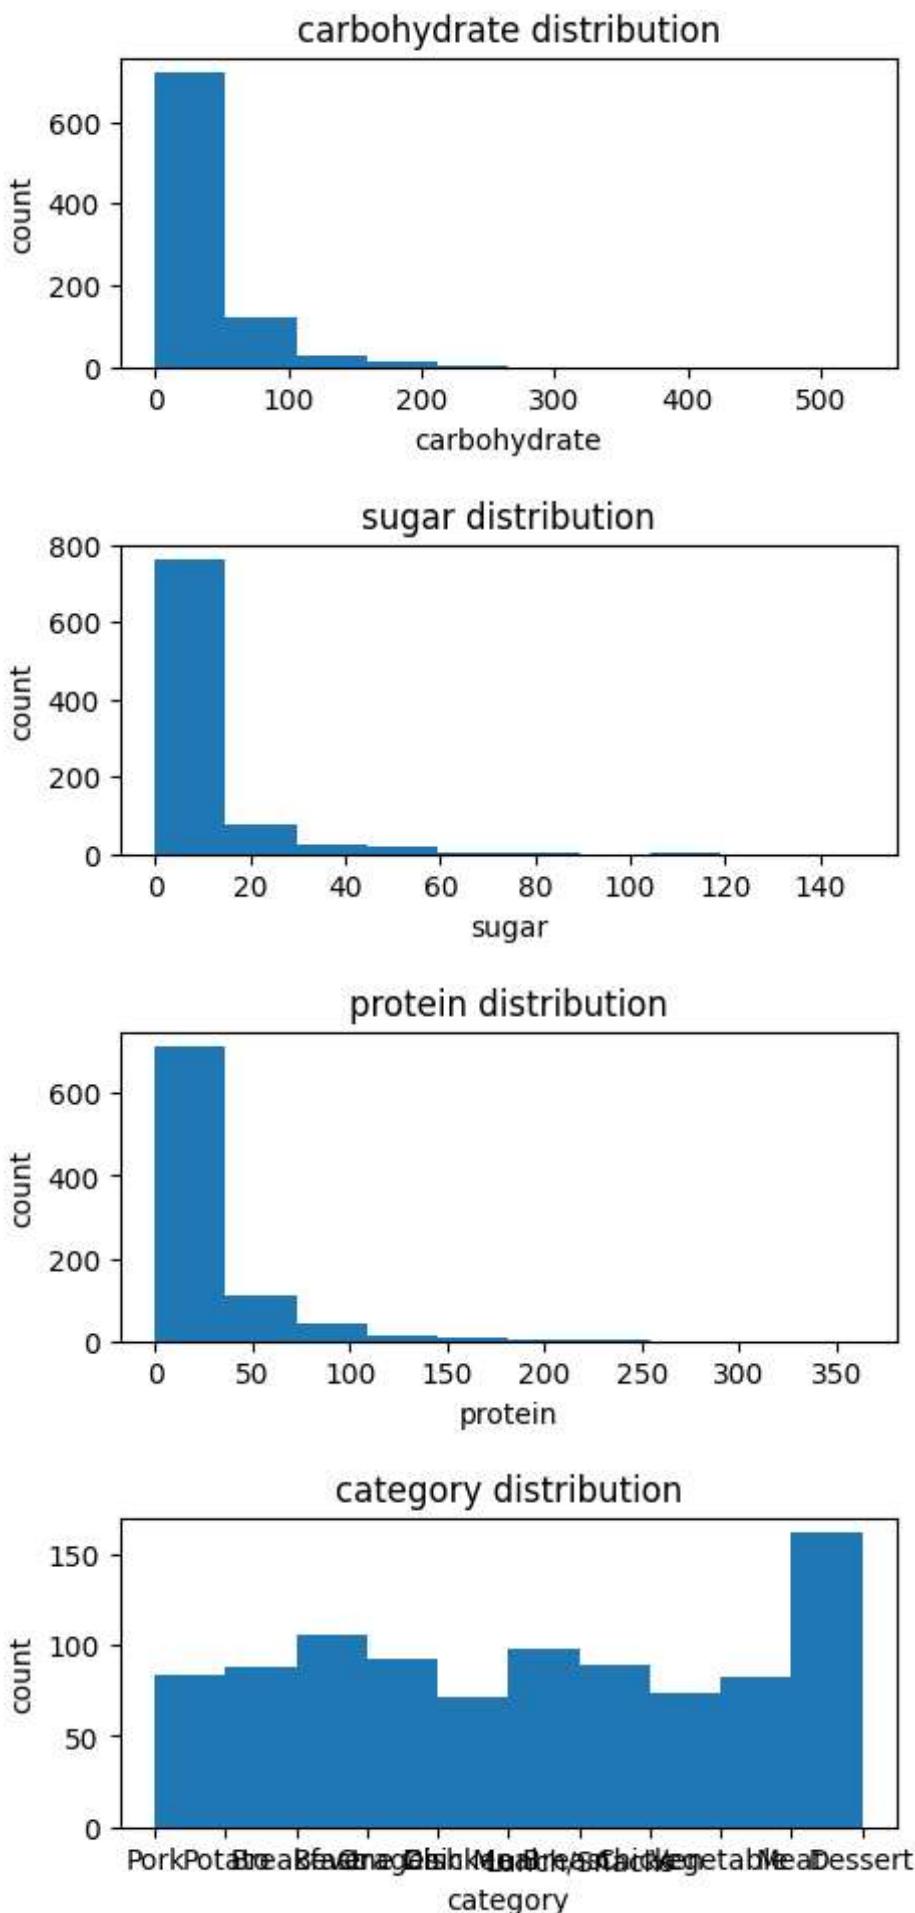
```
In [7]: df["category"].value_counts()
```

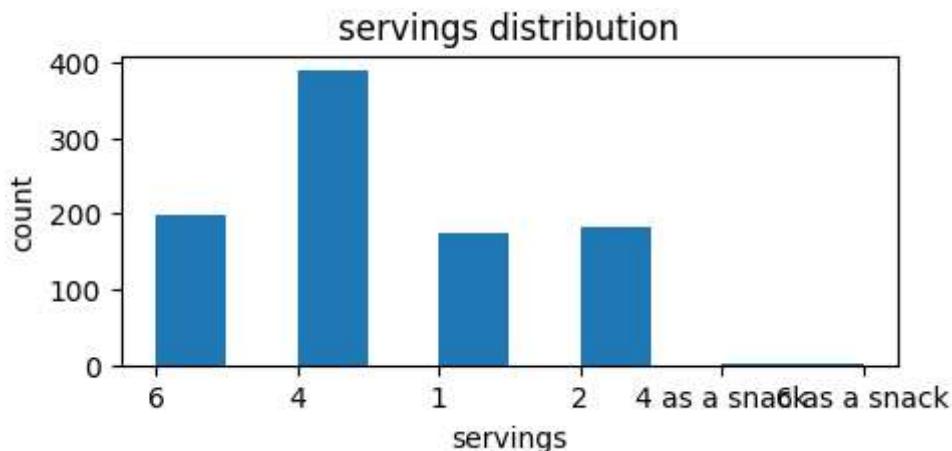
```
Out[7]: category
Breakfast      106
Chicken Breast 98
Beverages      92
Lunch/Snacks   89
Potato         88
Pork           84
Vegetable       83
Dessert         83
Meat            79
Chicken          74
One Dish Meal    71
Name: count, dtype: int64
```

```
In [8]: import matplotlib.pyplot as plt
```

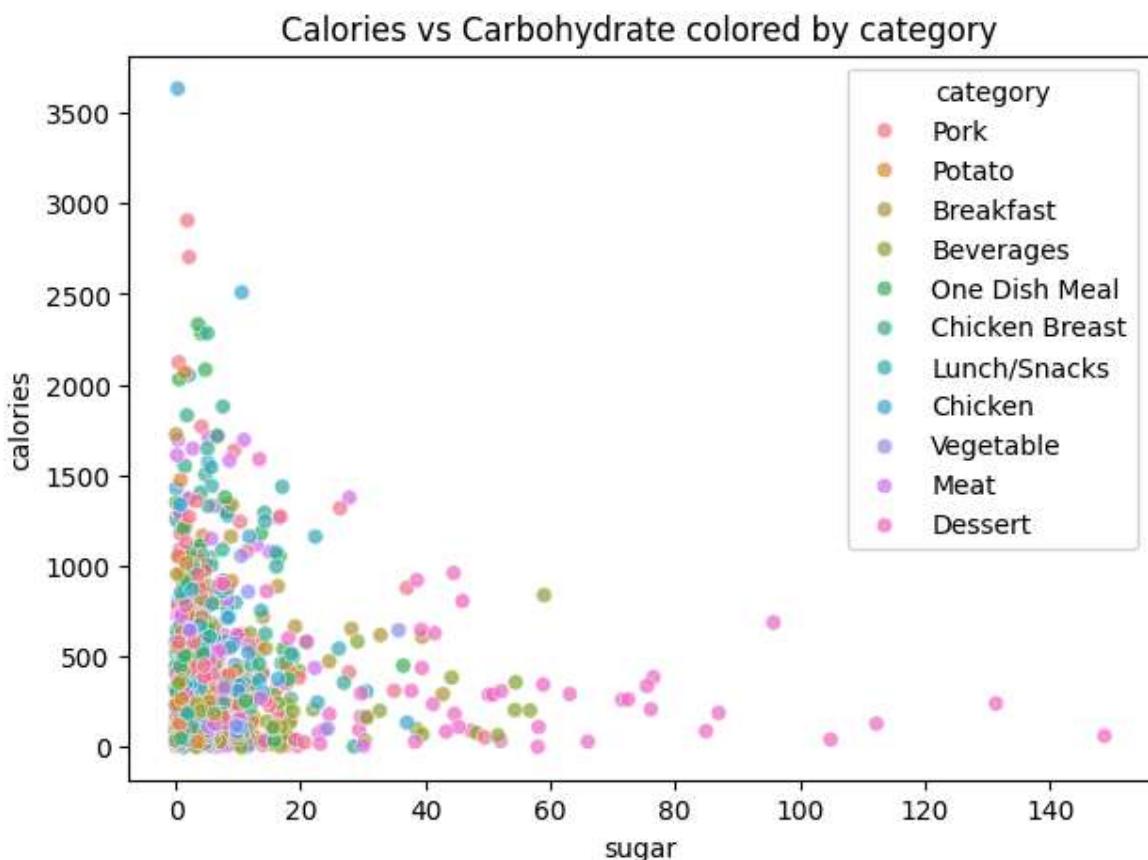
```
In [9]: for i in df.columns.drop("high_traffic"):
    plt.figure(figsize=(5, 2))
    plt.hist(df[i])
    plt.xlabel(i)
    plt.ylabel("count")
    plt.title(i+" distribution")
    plt.show()
```







```
In [11]: import seaborn as sns
plt.figure(figsize=(7, 5))
sns.scatterplot(
    data=df,
    x="sugar",
    y="calories",
    hue="category",
    alpha=0.7
)
plt.title("Calories vs Carbohydrate colored by category")
plt.show()
```



```
In [12]: df["servings"].value_counts()
```

```
Out[12]: servings
4           389
6           197
2           183
1           175
4 as a snack      2
6 as a snack      1
Name: count, dtype: int64
```

```
In [13]: #Impute
#for right-skewed columns, impute NAs with median
from sklearn.impute import SimpleImputer
num_cols = ["calories", "carbohydrate", "sugar", "protein"]

imputer = SimpleImputer(strategy="median")
df[num_cols] = imputer.fit_transform(df[num_cols])

#for high_traffic, impute NAs with Low
df["high_traffic"] = df["high_traffic"].fillna("Low")
```

```
In [14]: df.isna().sum()
```

```
Out[14]: recipe      0
calories     0
carbohydrate 0
sugar        0
protein      0
category     0
servings     0
high_traffic 0
dtype: int64
```

```
In [15]: df["servings"] = df["servings"].astype(str).str.extract(r"(\d+)").astype(int)
```

```
In [16]: df["servings"].value_counts()
```

```
Out[16]: servings
4    391
6    198
2    183
1    175
Name: count, dtype: int64
```

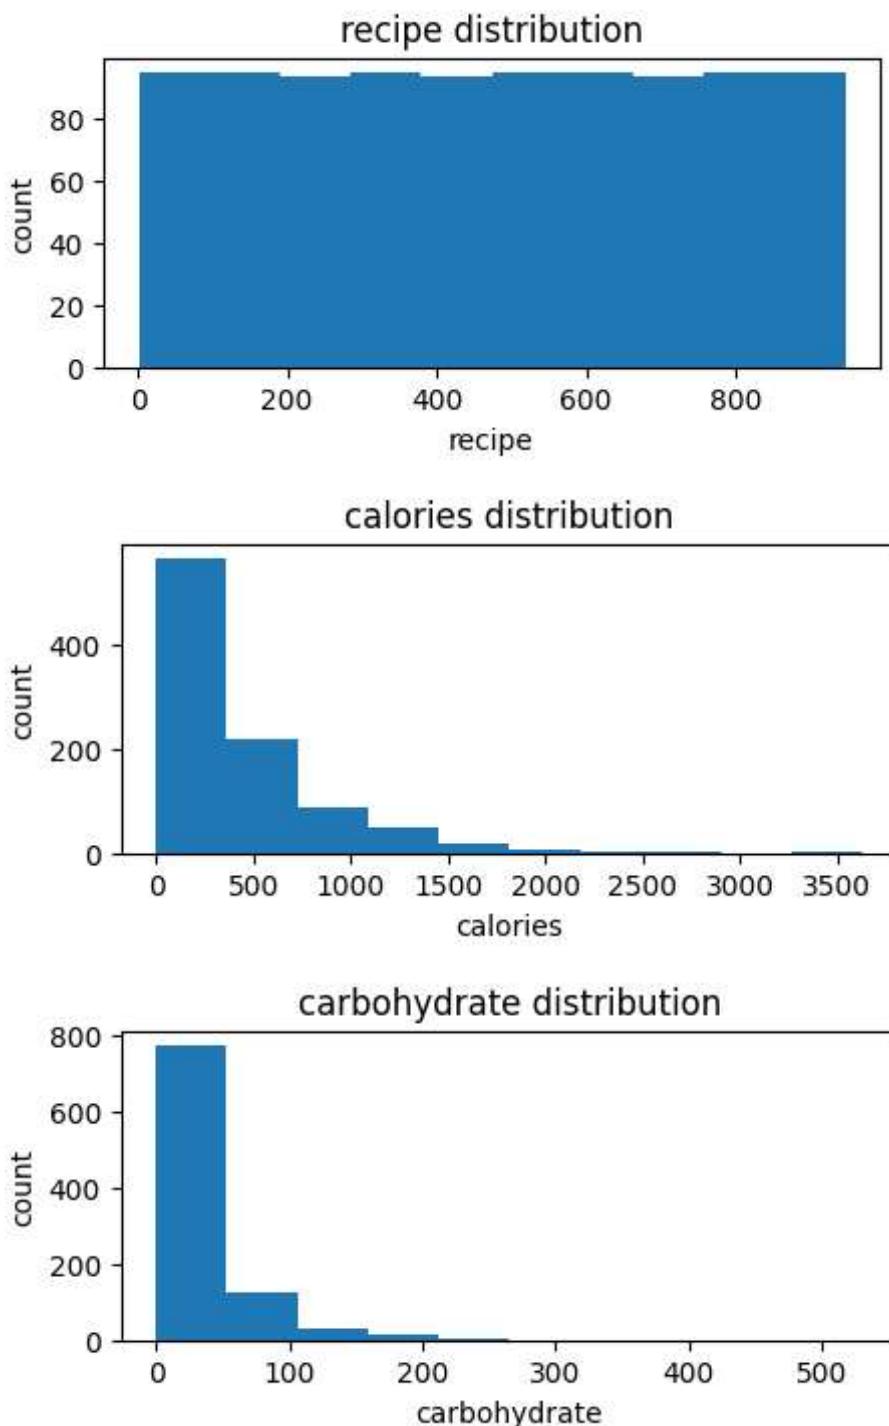
```
In [17]: df["category"] = df["category"].replace({"Chicken Breast": "Chicken"})
```

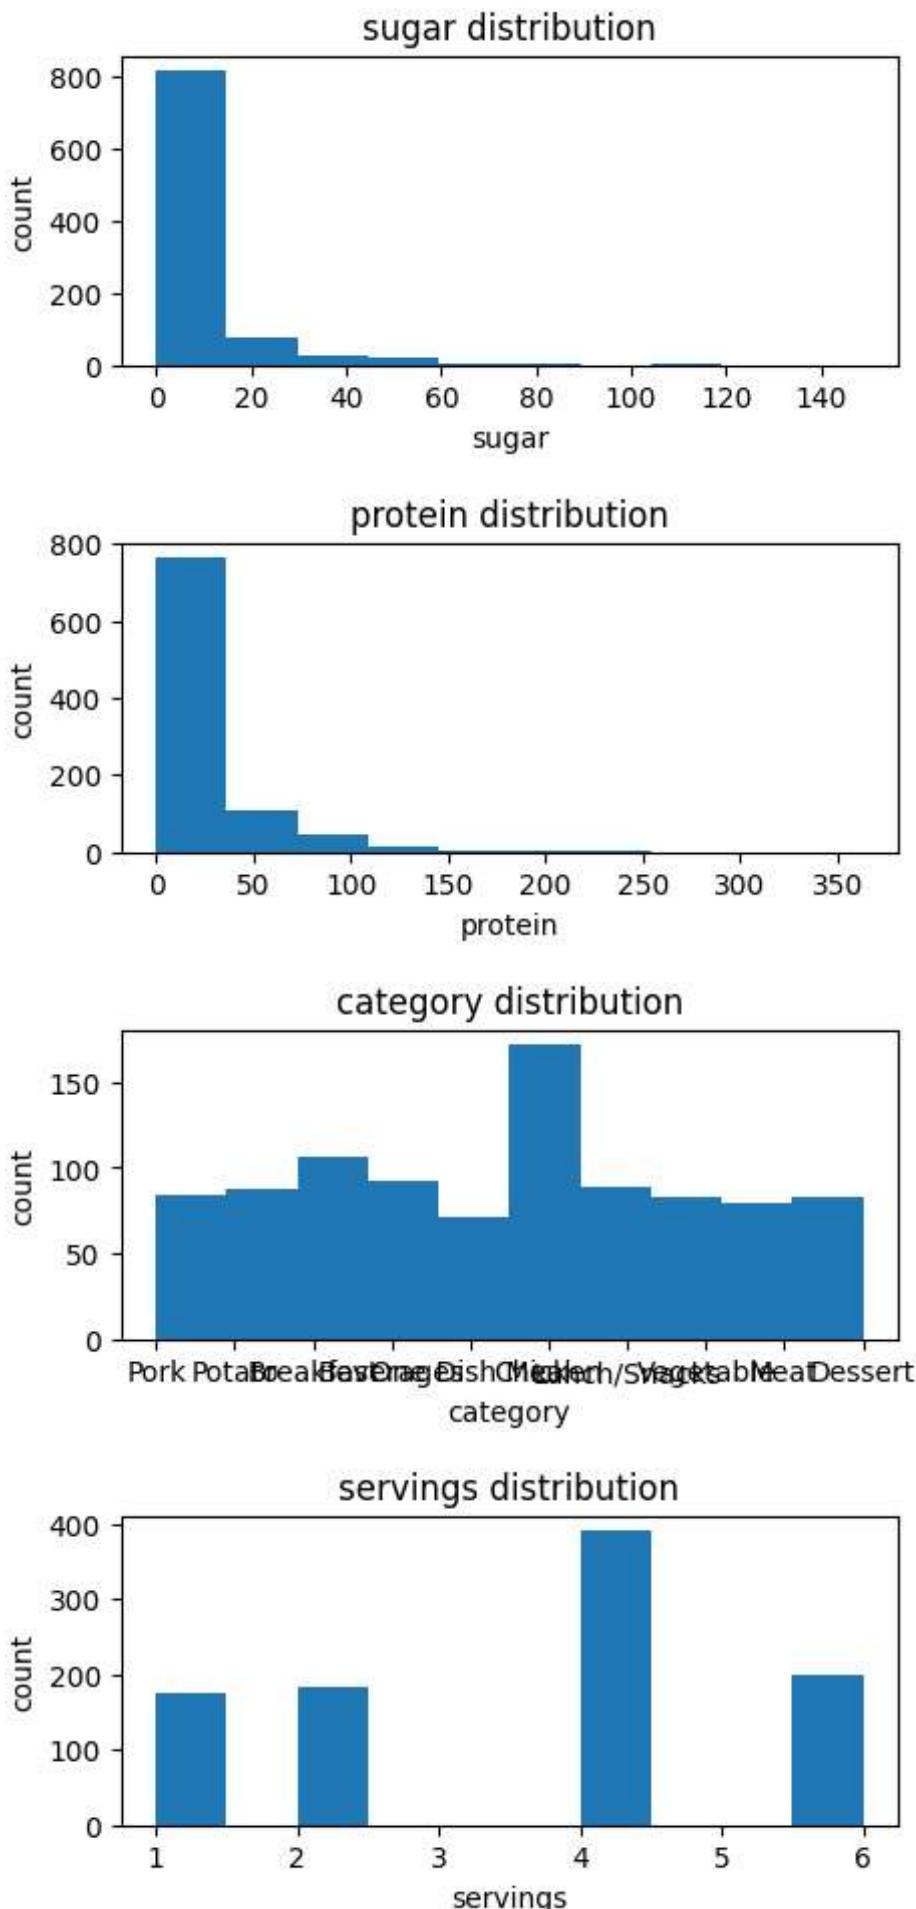
```
In [18]: for i in df.columns:
    print(df[i].dtypes)
```

```
int64
float64
float64
float64
float64
object
int32
object
```

```
In [19]: for i in df.columns.drop("high_traffic"):
    plt.figure(figsize=(5, 2))
```

```
plt.hist(df[i])
plt.xlabel(i)
plt.ylabel("count")
plt.title(i+" distribution")
plt.show()
```

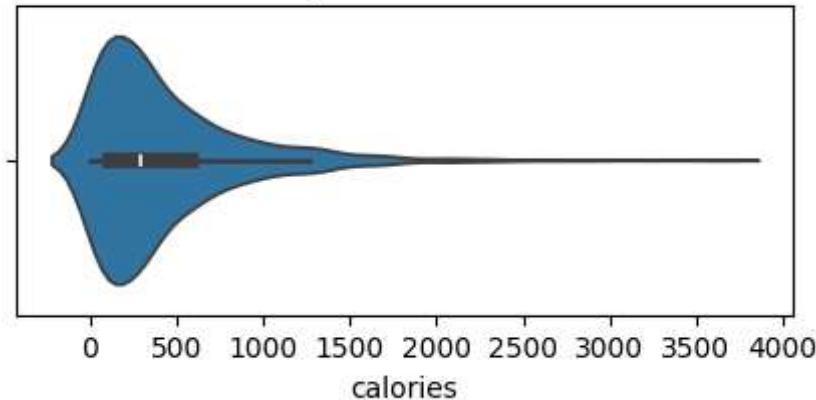




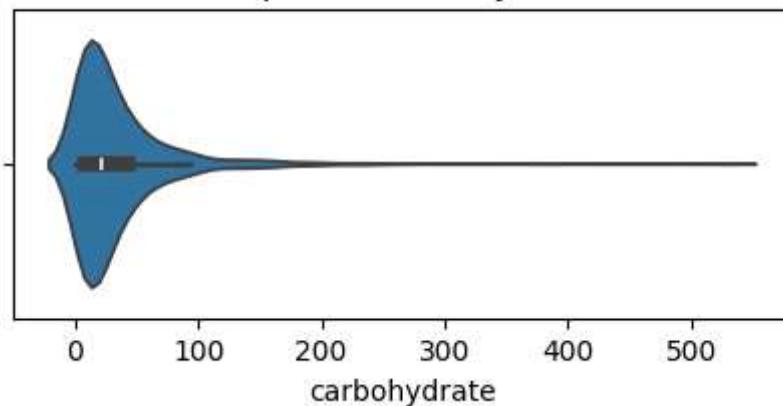
```
In [20]: for i in df.columns.drop(["high_traffic","recipe"]):
    plt.figure(figsize=(5,2))
```

```
sns.violinplot(x=df[i])
plt.title("Boxplot of " + i)
plt.show()
```

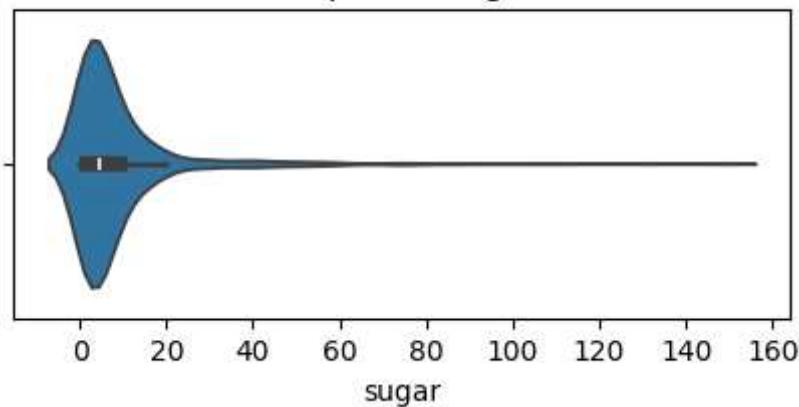
Boxplot of calories



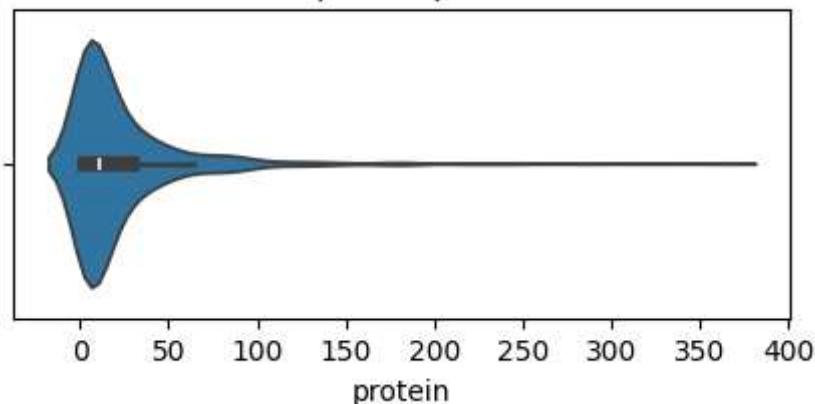
Boxplot of carbohydrate



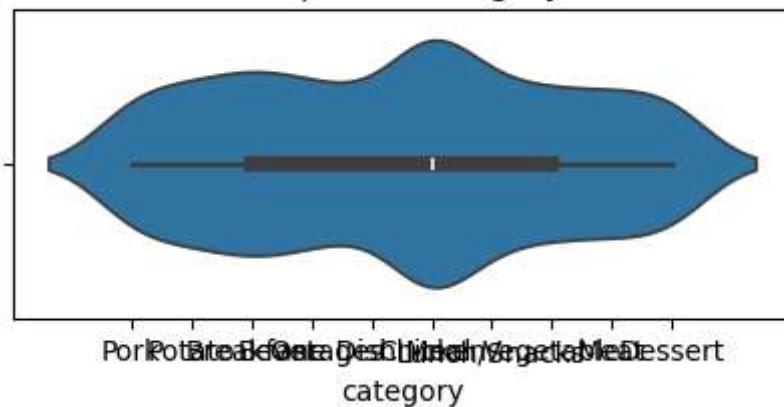
Boxplot of sugar



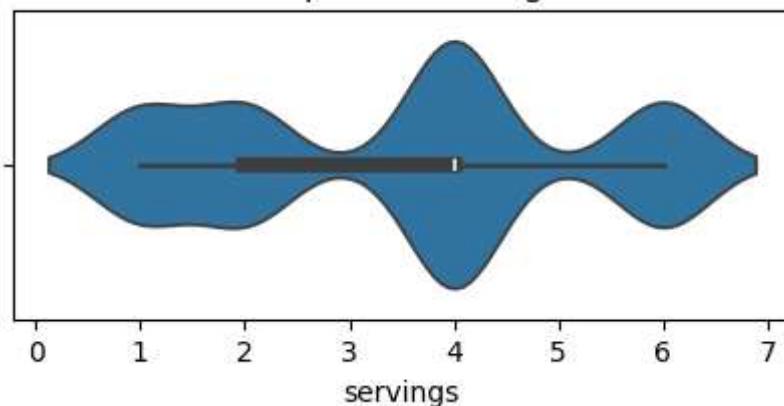
Boxplot of protein



Boxplot of category



Boxplot of servings



```
In [21]: from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder(handle_unknown="ignore", sparse_output=False)

encoded_cat = encoder.fit_transform(df[["category"]])

# 转回 DataFrame
encoded_df = pd.DataFrame(encoded_cat, columns=encoder.get_feature_names_out(["category"]))

# 合并回数据
df_encoded = pd.concat([df.drop(columns=["category"]), encoded_df], axis=1)

df_encoded["high_traffic"] = df_encoded["high_traffic"].map({"High": 1, "Low": 0})
X = df_encoded.drop(columns=["high_traffic", "recipe"])
y = df_encoded["high_traffic"]
```

```
In [22]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
In [23]: from sklearn.linear_model import LogisticRegression
logit = LogisticRegression(max_iter=50000)
logit.fit(X_train, y_train)
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
y_pred = logit.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Accuracy: 0.7736842105263158

	precision	recall	f1-score	support
0	0.69	0.79	0.73	75
1	0.85	0.77	0.80	115
accuracy			0.77	190
macro avg	0.77	0.78	0.77	190
weighted avg	0.78	0.77	0.78	190

Confusion Matrix:

```
[[59 16]
 [27 88]]
```

```
In [24]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

rf = RandomForestClassifier(
    n_estimators=300,
    random_state=42,
    class_weight='balanced',
)

rf.fit(X_train, y_train)

y_pred = rf.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 0.7315789473684211

	precision	recall	f1-score	support
0	0.64	0.72	0.68	75
1	0.80	0.74	0.77	115
accuracy			0.73	190
macro avg	0.72	0.73	0.72	190
weighted avg	0.74	0.73	0.73	190

```
[[54 21]
 [30 85]]
```

```
In [25]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

logit = LogisticRegression(max_iter=50000)
logit.fit(X_train, y_train)

y_proba = logit.predict_proba(X_test)[:, 1]

threshold = 0.43
y_pred_new = (y_proba >= threshold).astype(int)

print("Threshold:", threshold)
print("Accuracy:", accuracy_score(y_test, y_pred_new))
print("\nClassification Report:\n", classification_report(y_test, y_pred_new))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_new))
```

Threshold: 0.43

Accuracy: 0.7684210526315789

Classification Report:

	precision	recall	f1-score	support
0	0.70	0.72	0.71	75
1	0.81	0.80	0.81	115
accuracy			0.77	190
macro avg	0.76	0.76	0.76	190
weighted avg	0.77	0.77	0.77	190

Confusion Matrix:

```
[[54 21]
 [23 92]]
```